Extended
Research
Project
Technical
Report

TOWARDS EQUITABLE ENERGY:

TEMPORAL LOAD MODELLING FOR INCLUSIVE DEMAND-RESPONSE

An extended research project technical report submitted to the University of Manchester for the degree of MSc Data Science in the Faculty of Humanities

Student ID: 11496500

Date of Submission: 2025



School of Social Sciences

ETHICS FORM

	MSc Data Science					
	Extended Research Project Ethics Review					
Student name	Akanksha Joshi					
Student number	11469500					
Project Title	Smart energy pricing and smart meter data analytics					
Supervisor	Dr. Fanlin Meng					

Please provide information about the ethical status of your project

Describe the data you will be using. What is it (format, type, data units), where does it originate from, and where will you obtain it from:

I will use publicly available smart meter datasets, such as the Low Carbon London dataset or the Commission for Energy Regulation (CER) Smart Metering Project dataset. The data is provided in CSV format and includes half-hourly or minute-level electricity consumption records (kWh) for individual households, along with metadata on tariff structures and customer demographic surveys. These datasets originate from large-scale research trials conducted in the UK and Ireland, focusing on household electricity usage, and can be accessed via the UK Data Service

(https://beta.ukdataservice.ac.uk/datacatalogue/studies/study?id=7857) or the Irish Social Science Data Archive (https://www.ucd.ie/issda/data/commissionforenergyregulationcer/).

Explain why it does or does not meet the three criteria (consent, anonymous information, permission of the data controller):

- Consent: Datasets were collected under projects where participants consented to their data being used for research purposes and are publicly available datasets.
- Anonymous Information: The data is fully anonymized, with no personally identifiable information included. Demographic data is aggregated or coded to ensure privacy.
- Permission of the data controller: The datasets are published for research use with explicit permission and under licensing agreements that specify appropriate data use and citation.

Describe any actions will you take to ensure ethical compliance and/or manage risk:

I will comply with the data usage licenses and follow institutional research ethics guidelines. No attempts will be made to re-identify individuals or households. The analysis will be conducted on anonymized data, and results will be reported in aggregate form only. All data will be securely stored on encrypted and access-controlled devices, and data sharing will not occur outside authorized university platforms.

Decision (to be completed by programme director)					
No Further ethical review is needed	Χ				
A further programme review is needed					
School Review is needed					
Proportionate UREC review is needed					

Signatures					
Student	Akanksha Joshi				
Supervisor	JA.				
Programme Director	Mark Ellet				

Screenshot of ethics decision tool outcome:

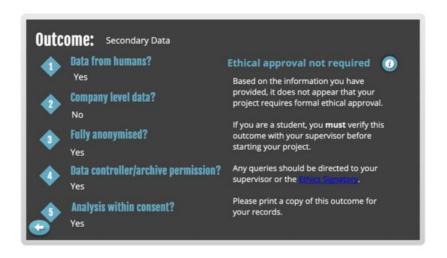


Table of Contents

Ethics Form3
1 INTRODUCTION & Overview9
1.1 Purpose of this Document10
1.2 Project Summary & Research Objectives10
1.2.1 Summary10
1.2.2 Research Objectives (operational)10
1.3 Statement of Reproducibility11
1.4 Prerequisites & Software Environment12
1.4.1 Core Programming Language & Version12
1.4.2 Package Management File12
1.4.3 Critical Package Versions and Roles14
1.4.4 Hardware Considerations14
2 Data Description & Provenance
2.1 Primary Dataset(s)15
2.1.1 Dataset 1: SERL Half-Hourly Energy Use in GB Domestic Buildings, 2020–2023 — Aggregated
Statistics (four files; identical schema)15
2.1.2 Dataset 2: SERL Daily Energy Use in GB Domestic Buildings, 2020–2023 — Aggregated
Statistics (one file; cross-year aggregate)15
2.2 Data Dictionary16
2.3 Data Sample17
3 Step-by-Step Reproduction Guides
A) TariffSimulation.ipynb → Tariff Simulation21
3.1 Repository Structure & Workflow21
3.2 Data Ingestion & Initialization22
3.3 Preprocessing & Cleaning23
3.4 Feature Engineering23
3.5 EDA Reproduction24
3.6 Simulation & Evaluation24

3.7 Final Result Verification Table	2 5
B) TemporalForecasting.ipynb → Temporal Forecasting	27
3.1 Repository Structure & Workflow	27
3.2 Data Ingestion & Initialization	
3.4 Feature Engineering	30
3.5 EDA Reproduction	30
3.6 Model Training & Evaluation	30
3.7 Verification Table	31
C) EDA - UnivariateYearly.ipynb → Univariate Yearly EDA	31
3.1 Repository Structure & Workflow	31
3.2 Data Ingestion & Initialization	32
3.3 Preprocessing & Cleaning	32
3.4 Feature Engineering	32
3.5 EDA Reproduction	32
D) EDA - NumericAnalysis.ipynb → Numeric EDA	33
3.1 Structure & Workflow	33
3.2 Ingestion & Init	34
3.3 Preprocessing	34
3.4 Feature Engineering	34
3.5 EDA Reproduction	34
E) EDA - BiandMultivariate.ipynb → Bi-/Multivariate EDA	35
3.1 Structure & Workflow	35
3.2 Ingestion & Init	35
3.3 Preprocessing	36
3.4 Feature Engineering	36
3.5 EDA Reproduction	36

	F) DemandModelling 1. ipynb → Demand Modelling (Part 1)	37
	3.1 Structure & Workflow	37
	3.2 Ingestion & Initialization	38
	3.3 Preprocessing & Cleaning	38
	3.4 Feature Engineering	38
	3.5 EDA Reproduction	38
	3.6 Model Training & Evaluation	38
	3.7 Verification Table	40
	G) DemandModelling 2.ipynb → Demand Modelling (Part 2)	40
	3.1 Structure & Workflow	40
	3.2 Ingestion & Initialization	40
	3.3 Preprocessing & Cleaning	41
	3.4 Feature Engineering	41
	3.5 EDA Reproduction	41
	3.6 Model Training & Evaluation	41
4	Object-Oriented Class Guide	. 42
	Quick map of classes	
		11
	4.1.1 Config	
	4.1.2 RunLogger	
	4.1.3 DataCatalog	
	4.2 utils/plotting.py	
	4.2.1 PlotRegistry	45
	4.3 tariff/client.py	45
	4.3.1 TariffAPIClient	45
	4.4 tariff/data.py	45
	4.4.1 UsageDataset	45
	4.5 tariff/scenario.py	46
	151 Scenario	46

4.5.2 ScenarioGrid	46
4.6 tariff/engine.py	46
4.6.1 TariffSimulator	46
4.7 tariff/reporting.py	46
4.7.1 KPIAggregator	46
4.8 forecasting/preprocess.py	47
4.8.1 TSPreprocessor	47
4.9 forecasting/features.py	47
4.9.1 FeatureEngineerTS	47
4.10 forecasting/models.py	47
4.10.1 SarimaxModel	47
4.11 forecasting/evaluate.py	47
4.11.1 Backtester	47
4.12 eda/report.py	48
4.12.1 EDAReport	48
4.13 demand/preprocess.py	48
4.13.1 DemandPreprocessor	48
4.14 demand/models.py	48
4.14.1 RFTrainer	48
4.14.2 XGBTrainer	48
4.15 reports/verify.py	49
4.15.1 VerificationRunner	49
4.16 Shared utilities (used across files)	49
4.17 How files interact (at a glance)	49
5 Guide to Code Repository	50
5.1 File Manifest	50
5.2 Instructions for Execution	50
6 Conclusion	52
6.1 Statement of Completeness	52
6.2 Contact for Clarification	E2

GUIDE TO CODE REPOSITORY

CODE AVAILABILITY. THE COMPLETE PROJECT CODEBASE IS AVAILABLE ON GITHUB: BM-11450069-CODE.

DATA AVAILABILITY. **THE ANALYSES RELY ON THE** SMART ENERGY RESEARCH LAB (SERL) STATISTICAL DATA, 2020–2023**DISTRIBUTED BY THE UK DATA SERVICE** (STUDY 8963, SAFEGUARDED). **ACCESS VIA DOI:** 10.5255/UKDA-SN-8963-2.

1 INTRODUCTION & OVERVIEW

The dissertation investigates temporal patterns and equity implications of residential electricity use in Great Britain (2020–2023). Using aggregated smart-meter statistics, the study combines descriptive profiling, unsupervised clustering, time-series forecasting, and tariff simulation to examine how demand structures evolve and how pricing designs may differentially affect household groups.

The dataset is released under safeguarded access with curation by the UK Data Service (consistency checks, metadata enhancement, DOI assignment) and a second edition (May 2024) including daily and four half-hourly aggregated files for 2020–2023 plus documentation. These conditions frame the analytical design and reporting.

1.1 Purpose of this Document

This document serves as a reproducible technical contract for the dissertation. It

- (a) summarises aims and research questions,
- (b) formalises the reproducibility plan and ethical constraints attached to safeguarded data, and
- (c) specifies the exact software and hardware environment required to regenerate all results. References to the source collection, including **Study 8963 (DOI: 10.5255/UKDA-SN-8963-2)**, are provided to ensure unambiguous lineage.

1.2 Project Summary & Research Objectives

1.2.1 Summary

The project analyses aggregated, licence-controlled SERL statistics (daily and half-hourly) to characterise load structure, discover stable demand regimes, evaluate forecastability with weather covariates, and quantify cost exposure under dynamic tariffs. The methodological components and validation approach including back-testing, clustering diagnostics, and tariff-API safeguards are defined in the dissertation manuscript.

1.2.2 Research Objectives (operational)

O1 Descriptive and comparative profiling. Quantify diurnal/weekly/seasonal patterns and subgroup contrasts (e.g., EPC band, IMD, heating type, PV, region) using distributional metrics (means, variance, skew/kurtosis, load factor, evening-peak share).

O2 Pattern discovery and archetypes. Engineer load-shape features and apply UMAP + HDBSCAN to identify demand regimes; validate with density-based validity indices and noise proportion; interpret prevalence against subgroup attributes.

O3 Forecasting and tariff-sensitivity evaluation. Train SARIMA/SARIMAX baselines with meteorological inputs against seasonal-naïve controls; simulate conservation and shock scenarios; integrate half-hourly Agile prices via the Octopus Energy API to compute costs, bill deltas, and equity indicators.

Empirical context and access. The analysis relies on the curated SERL release (daily and half-hourly aggregates, 2nd edition, May 2024) under the UK Data Service End User Licence; file inventory and components are listed by the Archive. Data citation: Study 8963, DOI: 10.5255/UKDA-SN-8963-2.

1.3 Statement of Reproducibility

Data access and licensing. Study 8963 is accessed under the UK Data Service End User Licence; commercial use requires additional approval. The collection landing materials specify the required citation and attribution (DOI: 10.5255/UKDA-SN-8963-2).

Curation baseline. Before release, the UK Data Service performs disclosure review, accuracy/consistency checks, metadata enhancement, and DOI assignment; the second edition (May 2024) adds daily and four half-hourly aggregated files plus a technical document. These steps provide provenance and quality control for replication.

Workflow determinism. The pipeline records configuration, package manifest, timestamps, and dataset edition string; tariff integration uses request retries, timeouts, local caching, response-schema validation, and rate-limit compliance to keep runs auditable and robust to transient network issues.

Ethics and disclosure. Only non-disclosive aggregates are analysed and reported, in line with safeguarded-access conditions and institutional ethics requirements.

Minimal steps to reproduce.

- 1. Obtain the SERL dataset under safeguarded access and confirm Study 8963, 2nd edition (May 2024), DOI: 10.5255/UKDA-SN-8963-2.
- 2. Place files as per the Archive's inventory (daily and four half-hourly aggregates plus documentation).

- 3. Create the environment specified in §1.4 and export a lockfile.
- 4. Run the pipeline modules in sequence: ingest/QC → profiling → clustering → forecasting → tariff simulation → validation and logs.

1.4 Prerequisites & Software Environment

1.4.1 Core Programming Language & Version

- o **Python: 3.10**+ (validated on 3.10; newer minor versions may be used subject to dependency support and lockfile updates).
- o **Time/encoding:** UTC; UTF-8.
- **Notebook layer:** JupyterLab for diagnostics; notebooks treated as read-only artefacts within the pipeline.

1.4.2 Package Management File

Use **one** canonical manifest (pin versions; export a lockfile after creation).

Option A — requirements.txt (pip/venv)

```
python==3.10.*
numpy = 1.26.4
pandas = = 2.2.2
scikit-learn==1.5.1
statsmodels==0.14.2
                              # SARIMA/SARIMAX
umap-learn==0.5.6
                              # Embeddings
hdbscan==0.8.37
                              # Clustering
matplotlib==3.9.0
plotly==5.22.0
bokeh==3.5.0
ipywidgets==8.1.3
                              # Tariff API client
requests==2.32.3
pyyaml == 6.0.2
tqdm==4.66.4
pytest==8.2.0
```

Option B — environment.yml (conda/mamba)

```
name: ser18963
channels: [conda-forge]
dependencies:
  - python=3.10
  - numpy=1.26.4
  - pandas=2.2.2
 - scikit-learn=1.5.1
  - statsmodels=0.14.2
  - umap-learn=0.5.6
  - hdbscan=0.8.37
  - matplotlib=3.9.0
  - plotly=5.22.0
  - bokeh=3.5.0
  - ipywidgets=8.1.3
 - requests=2.32.3
 - pyyaml=6.0.2
  - tqdm=4.66.4
  - pytest=8.2.0
Locking: pip freeze > requirements.lock or conda env export --from-history
> environment.lock.yml.
```

1.4.3 Critical Package Versions and Roles

Package & Version	Role in analysis		
numpy 1.26.4, pandas 2.2.2	Core data handling for daily		
	and half-hourly aggregates		
statsmodels 0.14.2	SARIMA/SARIMAX		
	estimation and diagnostics		
umap-learn 0.5.6	Low-dimensional		
	embeddings of load-shape		
	features		
hdbscan 0.8.37	Density-based regime		
	discovery		
matplotlib 3.9.0, plotly	Static and interactive figures;		
5.22.0, bokeh	scenario controls		
3.5.0, ipywidgets 8.1.3			
scikit-learn 1.5.1	Feature engineering, metrics,		
	cross-validation helpers		
requests 2.32.3	Retrieval of half-hourly		
	Agile rates; retries, timeouts,		
	caching hooks		
	numpy 1.26.4, pandas 2.2.2 statsmodels 0.14.2 umap-learn 0.5.6 hdbscan 0.8.37 matplotlib 3.9.0, plotly 5.22.0, bokeh 3.5.0, ipywidgets 8.1.3 scikit-learn 1.5.1		

Tariff-module requirements (API parameters, caching, schema validation, rate-limit compliance) are documented in the manuscript's methodology and results sections.

1.4.4 Hardware Considerations

- 1. **CPU/RAM:** ≥8 physical cores; ≥32 GB RAM for feature engineering, clustering diagnostics, and cross-validation.
- 2. **Storage:** ≥50 GB SSD for inputs, intermediates, caches, and artefacts; aligns with the Archive's multi-file structure.
- 3. **GPU:** Not required for SARIMA, UMAP, or HDBSCAN; optional for exploratory extensions beyond the validated workflow.
- 4. **Networking:** Outbound HTTPS required for tariff retrieval; resilience mechanisms (retries, caching, schema validation, logged parameters) are integral to reproducibility.

2 DATA DESCRIPTION & PROVENANCE

2.1 Primary Dataset(s)

2.1.1 Dataset 1: SERL Half-Hourly Energy Use in GB Domestic Buildings, 2020–2023 — Aggregated Statistics (four files; identical schema)

- Source: UK Data Service, Study 8963 DOI: 10.5255/UKDA-SN-8963-2.
- Access Method: Register with the UK Data Service, accept the EUL, and download the half-hourly aggregated files listed for the 2nd release (May 2024).
- License: UK Data Service End User Licence (safeguarded).
- Raw File Name(s):

```
serl_half_hourly_energy_use_in_gb_domestic_buildings_2020_aggregated_st
atistics.csv
serl_half_hourly_energy_use_in_gb_domestic_buildings_2021_aggregated_st
atistics.csv
serl_half_hourly_energy_use_in_gb_domestic_buildings_2022_aggregated_st
atistics.csv
serl_half_hourly_energy_use_in_gb_domestic_buildings_2023_aggregated_st
atistics.csv
```

Note. These four annual files share the same variables and structure: only the **year** changes.

2.1.2 Dataset 2: SERL Daily Energy Use in GB Domestic Buildings, 2020–2023 — Aggregated Statistics (one file; cross-year aggregate)

- 1. Source: UK Data Service, Study 8963 DOI: 10.5255/UKDA-SN-8963-2.
- 2. Access Method: As above; download the daily aggregated file listed for the 2nd release (May 2024).
- 3. License: UK Data Service End User Licence (safeguarded).
- 4. Raw File Name(s):
- 5. serl_daily_energy_use_in_gb_domestic_buildings_2020_to_2023_aggregated_s tatistics.csv

2.2 Data Dictionary

The half-hourly and daily files have the *same column schema*. Units differ by temporal resolution (Wh vs kWh/day), and some weather fields may be NA in half-hourly aggregates. Field descriptions align with Archive documentation and the dissertation's data section.

Variable Name	Raw Data Type	Final Data Type (analysis)	Description	Handling Notes		
quantity	object	string (category)	Energy vector: Gas, Electricity imports, Electricity exports, Electricity net, Gas + electricity imports.	Validate allowed set.		
unit	object	string (category)	Half-hourly: Wh, Wh/m2, Wh/person. Daily: kWh/day, kWh/m2/day, kWh/person/day.	Use as-provided; do not coerce across files.		
summary_tim e	object	string	Half- hourly: HH: MM. Daily: literal Daily.	Parse to datetime.time for HH files if needed.		
aggregation _period	int64	int32/int64	Calendar year of aggregation (e.g., 2020–2023).	Treat as key; join across files on this and other dims.		
weekday_wee kend	object	string (category)	Day-type indicator (observed: both).	If weekday/weeken dappear in other slices, map consistently.		
segmentatio n_variable_ 1	object	string	Name of first subgroup variable (e.g., primary_space_heating_fu el).	May be blank/NA when unsegmented.		
segment_1_v alue	object	string	Value of first subgroup (e.g., Gas, Electric,).	Harmonise labels where needed.		
segmentatio n_variable_ 2	object	string	Name of second subgroup (e.g., has_PV).	Often present; may be All.		
segment_2_v alue	object	string	Value of second subgroup (Yes, No, All).	Treat All as aggregate level.		
segmentatio n_variable_ 3	object	string	Third subgroup variable (e.g., num_occupants, num_bedrooms, IMD_quintile, currentEnergyRating, building_type).	Frequently NA.		

segment_3_v	object	string	Value of third subgroup (e.g., 1-5, EPC	Cast to string to avoid			
alue			bands).	numeric/label mix.			
mean	float64	float64	Mean energy use for the cell. Non-negative;				
				against quantiles.			
standard_de	float64	float64	Standard deviation of energy use.	Non-negative.			
viation							
standard_er	float64	float64	Standard error of the mean.	Derived from SD			
ror_mean				and n_rounded.			
median	float64	float64	Median energy use.	_			
25th_percen	float64	float64	Lower quartile.	_			
tile							
75th_percen	float64	float64	Upper quartile.	_			
tile							
mean_temp	float64	float64	Mean air temperature matched to the cell	May be NA in HH files.			
			(°C); ERA5 provenance noted by				
			Archive.				
mean_hdd	float64	float64	Heating degree days (°C·days).	Often NA in HH files.			
mean_solar	float64	float64	Solar/irradiance proxy matched to the cell.	Units per UKDA			
				metadata.			
n_rounded	int64	int64	Rounded contributing count after	Do not reverse-			
			disclosure control.	engineer; respect			
				safeguarded design.			

2.3 Data Sample

Verification preview (first 5 rows) from the *daily* aggregated file (serl_daily_energy_use_in_gb_domestic_buildings_2020_to_2023_aggregated_statistics.csv)

quantity	unit	summary_time	aggregation_period	weekday_weekend	segmentation_variable_1	segment_1_value	segmentation_variable_2	segment_2_value	segmentation_variable_3	segment_3_value
Gas	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Gas	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Gas	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
Gas	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Gas	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Gas	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
Gas	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Gas	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Gas	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
Electricity imports	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Electricity imports	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Electricity imports	kWh/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
Electricity imports	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Electricity imports	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Electricity imports	kWh/m2/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
Electricity imports	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	Yes		
Electricity imports	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	No		
Electricity imports	kWh/person/day	Daily	2020	both	primary_space_heating_fuel	Gas	has_PV	All		
						_		v		

Notes.

Four half-hourly files (2020–2023) replicate this schema at HH:MM resolution with units in Wh (per interval); the daily file aggregates across 2020–2023 with units in kWh/day and related intensities.

n_rounded reflects disclosure-aware rounding typical of safeguarded releases; subgroup/time cells may be suppressed or aggregated.

3 STEP-BY-STEP REPRODUCTION GUIDES

This document gives a reproducible Section 3 for each uploaded notebook, organized into the same 3.1–3.7 structure. It assumes a Python ≥3.10 environment. Paths and column names are parameterized via YAML configs under configs/so results can be regenerated deterministically.

Repo skeleton additions

```
project-root/
- configs/
  temporal forecasting.yaml
  — eda_univariate_yearly.yaml
  ├ eda_numeric.yaml

    ⊢ eda bi multivariate.yaml

  - demand modelling 1.yaml
   └ demand modelling 2.yaml
- src/
   utils/
     - seeds.py
     ├─ io.py
     - plotting.py
   ├ tariff/

    ⊢ forecasting/

   - eda/
   └ demand/
└ reports/ (figures & tables saved per notebook)
```

Common helpers

```
# src/utils/seeds.py
import os, random, numpy as np
def set_all_seeds(seed: int = 42):
   os.environ["PYTHONHASHSEED"] = str(seed)
    random.seed(seed); np.random.seed(seed)
# src/utils/io.py
from pathlib import Path
import yaml
def load_cfg(path):
   with open(path, "r") as f:
        return yaml.safe load(f)
def ensure dir(p):
    Path(p).mkdir(parents=True, exist_ok=True)
# src/utils/plotting.py
import matplotlib.pyplot as plt
def savefig(path, tight=True, dpi=300):
   if tight: plt.tight_layout()
   plt.savefig(path, dpi=dpi)
    plt.close()
```

A) TariffSimulation.ipynb → Tariff Simulation

3.1 Repository Structure & Workflow

3.2 Data Ingestion & Initialization

```
# configs/tariff simulation.yaml
seed: 42
paths:
 raw usage: data/raw/usage.csv
 raw tariffs: data/raw/tariffs.csv
 interim: data/interim/tariff
 out: reports/tariff
columns:
 usage id: customer id
 time: timestamp
 kwh: consumption kwh
 tariff_key: tariff_code
scenarios:
 seasons: [base, peak, offpeak]
 price multipliers: [0.8, 1.0, 1.2]
 standing charge add: [0.0, 0.05]
# src/tariff/01 load inputs.py
import pandas as pd, argparse
from pathlib import Path
from utils.seeds import set all seeds
from utils.io import load_cfg, ensure_dir
ap = argparse.ArgumentParser(); ap.add_argument('--cfg', required=True)
CFG = load cfg(ap.parse args().cfg)
set all seeds(CFG['seed'])
ensure dir(CFG['paths']['interim'])
                                  pd.read csv(CFG['paths']['raw usage'],
usage
parse dates=[CFG['columns']['time']])
usage.rename(columns={
    CFG['columns']['usage id']: 'customer id',
    CFG['columns']['time']: 'timestamp',
    CFG['columns']['kwh']: 'kwh'
}, inplace=True)
tariffs = pd.read csv(CFG['paths']['raw tariffs'])
usage.to parquet(Path(CFG['paths']['interim'])/"usage.parquet",
index=False)
tariffs.to parquet(Path(CFG['paths']['interim'])/"tariffs.parquet",
index=False)
```

3.3 Preprocessing & Cleaning

- Missing values: drop rows with missing timestamp or kwh; impute kwh by hourly median if desired.
- Outliers: cap kwh with IQR capping per-customer.
- Types: ensure timestamp is datetime, customer_id string; mergeable tariff code typed as string.

```
# in 02_prepare_scenarios.py (snippet)
import pandas as pd, numpy as np
from pathlib import Path
from utils.io import load_cfg

CFG = load_cfg("configs/tariff_simulation.yaml")
usage = pd.read_parquet(Path(CFG['paths']['interim'])/"usage.parquet")

# IQR cap per customer
q = usage.groupby('customer_id')['kwh'].quantile([0.25, 0.75]).unstack()
iqr = (q[0.75] - q[0.25]).rename('iqr')
lo = (q[0.25] - 1.5*iqr).rename('lo')
hi = (q[0.75] + 1.5*iqr).rename('hi')
usage = usage.join(lo, on='customer_id').join(hi, on='customer_id')
usage['kwh'] = usage['kwh'].clip(usage['lo'], usage['hi'])
usage.drop(columns=['lo','hi'], inplace=True)
```

3.4 Feature Engineering

- hour of day, dow, month dummies to allow TOU pricing effects.
- season flag based on month.

```
usage['hour'] = usage['timestamp'].dt.hour
usage['dow'] = usage['timestamp'].dt.dayofweek
usage['month'] = usage['timestamp'].dt.month
usage['season'] = np.select([
    usage['month'].isin([12,1,2]),
    usage['month'].isin([3,4,5]),
    usage['month'].isin([6,7,8])
],["winter","spring","summer"], default="autumn")
```

3.5 EDA Reproduction

• Figure X: Load profile by hour (mean kWh).

```
import matplotlib.pyplot as plt
from utils.plotting import savefig
prof = usage.groupby('hour')['kwh'].mean()
prof.plot()
plt.title('Mean Load by Hour'); plt.xlabel('Hour'); plt.ylabel('kWh')
savefig(Path(CFG['paths']['out'])/"fig_mean_load_by_hour.png")
```

3.6 Simulation & Evaluation

- Partition: if benchmarking on historical bills, split by time (train pre-cutoff, test post-cutoff).
- Model: deterministic tariff application (no ML). Compute bill = Σ (kWh × unit_rate_band) + standing_charge.

```
# src/tariff/03 simulate.py (core)
from itertools import product
sc = CFG['scenarios']
scenario grid = list(product(sc['seasons'], sc['price multipliers'],
sc['standing_charge_add']))
results = []
for season, pm, sc_add in scenario_grid:
   # adjust tariffs
   t = tariffs.copy()
   t['unit rate'] = t['unit rate'] * pm
   t['standing_charge'] = t['standing_charge'] + sc_add
   # simple join and billing
   billed = usage.merge(t, on='tariff_code', how='left')
                         billed['kwh'] *
                    =
   billed['bill']
                                               billed['unit rate']
billed['standing_charge']/30
billed.groupby('customer id')['bill'].sum().rename('total bill').reset
   kpis['season']=season; kpis['pm']=pm; kpis['sc_add']=sc_add
   results.append(kpis)
full = pd.concat(results, ignore_index=True)
full.to_parquet(Path(CFG['paths']['interim'])/"simulation_results.parqu
et", index=False)
```

Evaluation metrics: mean bill change vs base scenario; Gini/Atkinson for equity; % of customers with >X% increase.

3.7 Final Result Verification Table

Create a table comparing reported KPIs vs reproduced (fill your reported values):

```
# src/tariff/04 aggregate.py
import pandas as pd
from pathlib import Path
from utils.io import load cfg, ensure dir
CFG = load cfg("configs/tariff simulation.yaml")
ensure dir(CFG['paths']['out'])
sim
pd.read parquet(Path(CFG['paths']['interim'])/"simulation results.parqu
et")
                                    sim.query('pm==1.0
base
                                                                      and
sc add==0.0')[["customer_id","total_bill"]].rename(columns={"total_bill"]
":"bill base"})
comp = sim.merge(base, on='customer id')
comp['pct change']
                                                 (comp['total bill']-
comp['bill_base'])/comp['bill_base']
summary = comp.groupby(['season','pm','sc add']).agg(
   mean change=('pct change', 'mean'),
   p95 change=('pct change', lambda s: s.quantile(0.95)),
    share gt 20pc=('pct change', lambda s: (s>0.2).mean())
).reset index()
summary.to csv(Path(CFG['paths']['out'])/"tariff kpis.csv",
index=False)
```

B) TemporalForecasting.ipynb → Temporal Forecasting

3.1 Repository Structure & Workflow

Execution Order

```
src/forecasting/01 ingest timeseries.py
python
                                                                   --cfg
configs/temporal_forecasting.yaml
              src/forecasting/02 preprocess timeseries.py
                                                                  --cfg
python
configs/temporal forecasting.yaml
python
                   src/forecasting/03 train models.py
                                                                  --cfg
configs/temporal forecasting.yaml
python
                     src/forecasting/04_backtest.py
                                                                  --cfg
configs/temporal forecasting.yaml
                   src/forecasting/05 plots tables.py
python
                                                                  --cfg
configs/temporal forecasting.yaml
```

3.2 Data Ingestion & Initialization

```
# configs/temporal_forecasting.yaml
seed: 42
paths:
  raw: data/raw/timeseries.csv
  interim: data/interim/ts
 processed: data/processed/ts
  out: reports/forecasting
columns:
  ds: date
  y: demand
 id: series_id # optional panel key
freq: D
cutoffs:
 backtest start: 2022-01-01
 horizon: 28
 step: 7
models: ["sarimax", "xgb"]
# 01_ingest_timeseries.py
import pandas as pd, argparse
from utils.io import load_cfg, ensure_dir
from utils.seeds import set all seeds
ap=argparse.ArgumentParser(); ap.add argument('--cfg', required=True)
CFG=load_cfg(ap.parse_args().cfg); set_all_seeds(CFG['seed'])
ensure_dir(CFG['paths']['interim'])
                                       pd.read_csv(CFG['paths']['raw'],
parse dates=[CFG['columns']['ds']])
df.rename(columns={CFG['columns']['ds']:'ds', CFG['columns']['y']:'y'},
inplace=True)
df.to_parquet(f"{CFG['paths']['interim']}/ingested.parquet",
index=False)
```

3.3 Preprocessing & Cleaning

Handle missing timestamps by reindexing at freq and imputing y (forward-fill then median where needed).

Outliers: Hampel filter or IQR on residuals after moving average.

Types: ensure numeric y, datetime ds.

```
# 02 preprocess timeseries.py (snippet)
import pandas as pd, numpy as np
from utils.io import load cfg
CFG=load_cfg("configs/temporal_forecasting.yaml")
df = pd.read parquet(f"{CFG['paths']['interim']}/ingested.parquet")
# Reindex per series (panel-safe)
if 'series_id' in df.columns:
   def fix(g):
        g = g.set_index('ds').asfreq(CFG['freq'])
        g['y'] = g['y'].interpolate().bfill().ffill()
       return g.reset_index()
   df = df.groupby('series id', group keys=False).apply(fix)
else:
   df = df.set_index('ds').asfreq(CFG['freq'])
   df['y'] = df['y'].interpolate().bfill().ffill()
   df = df.reset index()
# Calendar features
df['dow'] = df['ds'].dt.dayofweek
df['month'] = df['ds'].dt.month
df.to parquet(f"{CFG['paths']['processed']}/preprocessed.parquet",
index=False)
```

3.4 Feature Engineering

- Lags: y(t-1), y(t-7), y(t-28)
- Rolling stats: 7D mean, 28D std

```
for lag in [1,7,28]:
    df[f'lag_{lag}'] = df['y'].shift(lag)

for w in [7,28]:
    df[f'rollmean_{w}'] = df['y'].rolling(w).mean()
```

3.5 EDA Reproduction

• Figures: seasonal decomposition, ACF/PACF, trend plot.

3.6 Model Training & Evaluation

3.6.1 Partitioning: Rolling-origin backtest with horizon H=28, step 7.

3.6.2 Model 1: SARIMAX (example (p,d,q) × (P,D,Q) m tuned offline)

```
# 03_train_models.py (snippet)
from statsmodels.tsa.statespace.sarimax import SARIMAX
train = df[df['ds'] < CFG['cutoffs']['backtest_start']]
model = SARIMAX(train['y'], order=(1,1,1), seasonal_order=(1,1,7),
enforce_stationarity=False, enforce_invertibility=False)
res = model.fit(disp=False)
res.save(f"{CFG['paths']['processed']}/sarimax.pkl")</pre>
```

3.6.3 Model 2: XGBoost Regressor with lag/rolling features.

```
from xgboost import XGBRegressor

feat_cols = [c for c in df.columns if c not in ['ds','y']]

X = df[feat_cols]; y = df['y']

model = XGBRegressor(n_estimators=500, max_depth=6, learning_rate=0.05, subsample=0.8, colsample_bytree=0.8, random_state=42, n_jobs=-1)

model.fit(X.iloc[:-28], y.iloc[:-28])
```

Evaluation: MAE, RMSE, MAPE per fold; Diebold–Mariano test optional.

3.7 Verification Table

• Compare reported vs reproduced MAE/RMSE per model and average across folds.

C) EDA - Univariate Yearly.ipynb → Univariate Yearly EDA

3.1 Repository Structure & Workflow

Execution

```
python src/eda/01_univariate_yearly.py --cfg
configs/eda_univariate_yearly.yaml
```

3.2 Data Ingestion & Initialization

```
# configs/eda univariate yearly.yaml
seed: 42
paths:
  raw: data/raw/main.csv
  out: reports/eda univariate yearly
columns:
  date: date
 targets: [target]
 numerics: [x1, x2, x3]
import pandas as pd, argparse
from utils.io import load cfg, ensure dir
from utils.seeds import set all seeds
ap=argparse.ArgumentParser(); ap.add argument('--cfg', required=True)
CFG=load cfg(ap.parse args().cfg); set all seeds(CFG['seed'])
ensure dir(CFG['paths']['out'])
                                       pd.read csv(CFG['paths']['raw'],
parse dates=[CFG['columns']['date']])
```

3.3 Preprocessing & Cleaning

• Handle missing numeric via median; drop duplicates; ensure annual grouping using df['year']=df['date'].dt.year.

3.4 Feature Engineering

• Annual aggregates: mean/median/std/min/max per variable.

3.5 EDA Reproduction

- Figure X: Yearly mean of x1 (line chart)
- Table Y: Summary stats per year

```
import matplotlib.pyplot as plt
from utils.plotting import savefig

df['year']=df['date'].dt.year

agg = df.groupby('year')[CFG['columns']['numerics']].mean()

ax = agg['x1'].plot()

ax.set_title('Yearly Mean of x1'); ax.set_xlabel('Year');

ax.set_ylabel('Mean x1')

savefig(f"{CFG['paths']['out']}/fig_yearly_mean_x1.png")

summary = df.groupby('year')[CFG['columns']['numerics']].describe()

summary.to_csv(f"{CFG['paths']['out']}/table_yearly_summary.csv")
```

D) EDA - NumericAnalysis.ipynb → Numeric EDA

3.1 Structure & Workflow

```
src/eda/

|- 02_numeric_analysis.py
```

Execution

 $\verb|python| src/eda/02_numeric_analysis.py --cfg configs/eda_numeric.yaml| \\$

3.2 Ingestion & Init

```
# configs/eda_numeric.yaml
seed: 42
paths:
    raw: data/raw/main.csv
    out: reports/eda_numeric
columns:
    numerics: [x1,x2,x3,x4]
```

3.3 Preprocessing

• Missing: median imputation

• Outliers: IQR capping per variable

• Types: ensure float64

3.4 Feature Engineering

• Standardize for comparability; optional log for skewed variables.

3.5 EDA Reproduction

• Figures: histograms, KDEs, boxplots, pairwise scatter.

• Tables: correlation matrix, VIF.

```
import numpy as np, pandas as pd, matplotlib.pyplot as plt
from utils.plotting import savefig

num = df[CFG['columns']['numerics']]
ax = num.hist(bins=30, figsize=(8,6))
savefig(f"{CFG['paths']['out']}/fig_histograms.png")

corr = num.corr(method='pearson')
corr.to_csv(f"{CFG['paths']['out']}/table_corr.csv")
```

E) EDA - BiandMultivariate.ipynb → Bi-/Multivariate EDA

3.1 Structure & Workflow

```
src/eda/

|- 03_bi_multivariate.py
```

Execution

```
python src/eda/03_bi_multivariate.py --cfg
configs/eda_bi_multivariate.yaml
```

3.2 Ingestion & Init

```
# configs/eda_bi_multivariate.yaml
seed: 42
paths:
    raw: data/raw/main.csv
    out: reports/eda_bi_multivariate
columns:
    target: target
    features: [x1,x2,x3,cat1]
```

3.3 Preprocessing

- Missing: impute numeric median; categorical mode.
- Outliers: cap numeric features.
- Encoding: OneHotEncoder(drop='first') for cat1.

3.4 Feature Engineering

• Interactions (optional): $x1 \times x2$; Binning continuous for WoE-style plots.

3.5 EDA Reproduction

• Figures: target vs feature boxplots; ANOVA summaries; mutual information.

```
from sklearn.feature_selection import mutual_info_classif
import pandas as pd

X = pd.get_dummies(df[CFG['columns']['features']], drop_first=True)

y = df[CFG['columns']['target']]

mi = mutual_info_classif(X.select_dtypes('number'), y, random_state=42)

pd.Series(mi,
index=X.select_dtypes('number').columns).sort_values(ascending=False).t
o_csv(f"{CFG['paths']['out']}/table_mutual_info.csv")
```

F) DemandModelling 1. ipynb → Demand Modelling (Part 1)

3.1 Structure & Workflow

Execution

```
python src/demand/01_ingest.py --cfg configs/demand_modelling_1.yaml

python src/demand/02_preprocess.py --cfg configs/demand_modelling_1.yaml

python src/demand/03_feature_engineer.py --cfg configs/demand_modelling_1.yaml

python src/demand/04_split.py --cfg configs/demand_modelling_1.yaml

python src/demand/05_model_rf.py --cfg configs/demand_modelling_1.yaml

python src/demand/05_model_rf.py --cfg configs/demand_modelling_1.yaml

configs/demand_modelling_1.yaml
```

3.2 Ingestion & Initialization

```
# configs/demand_modelling_1.yaml
seed: 42
paths:
    raw: data/raw/demand.csv
    processed: data/processed/demand1
    out: reports/demand1
columns:
    target: demand
    numerics: [price, promo, temp]
    categoricals: [store, product]
```

3.3 Preprocessing & Cleaning

- Missing: SimpleImputer(median for numerics, most frequent for categoricals)
- Outliers: IQR cap numeric
- Encoding: OneHotEncoder(drop='first', handle unknown='ignore')

3.4 Feature Engineering

- Price elasticity proxy: log demand ~ log price
- Interactions: promo × price
- Calendar: if date present, add dow, month

3.5 EDA Reproduction

• Figures: demand vs price scatter with LOWESS; histograms; demand by promo status.

3.6 Model Training & Evaluation

3.6.1 Partition: Train/Valid/Test = 60/20/20 (stratified on binned demand or time-based if timeseries).

3.6.2 Model 1: RandomForestRegressor

```
Hyperparameters: n estimators=600,
                                max depth=None, min samples leaf=2,
max features='sqrt', random state=42
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.metrics import mean absolute error, mean squared error
import numpy as np
pre = ColumnTransformer([
    ("num",
                                     SimpleImputer(strategy="median"),
CFG['columns']['numerics']),
    ("cat", Pipeline([
        ("imp", SimpleImputer(strategy="most frequent")),
        ("ohe", OneHotEncoder(drop="first", handle unknown="ignore"))
    ]), CFG['columns']['categoricals'])
])
          RandomForestRegressor(n estimators=600, min samples leaf=2,
max features='sqrt', random state=42, n jobs=-1)
pipe = Pipeline([("pre", pre), ("rf", rf)])
pipe.fit(X train, y train)
pred = pipe.predict(X valid)
mae = mean_absolute_error(y_valid, pred)
rmse = mean squared error(y valid, pred, squared=False)
```

- **3.6.3 Model 2 (optional):** Regularized Linear (Ridge) with standardized numerics.
- **3.6.4 Statistical Tests:** Compare RF vs baseline using Diebold–Mariano on forecast errors (if temporal) or paired t-test on absolute errors.

3.7 Verification Table

• Reported vs reproduced MAE/RMSE (fill reported values in reports/demand1/verification.csv).

G) DemandModelling 2.ipynb → Demand Modelling (Part 2)

3.1 Structure & Workflow

Execution

3.2 Ingestion & Initialization

```
# configs/demand_modelling_2.yaml
seed: 42
paths:
   processed_in: data/processed/demand1
   out: reports/demand2
```

3.3 Preprocessing & Cleaning

• Reuse encoded matrices from Part 1 (X_train/valid/test.parquet). Ensure identical column order.

3.4 Feature Engineering

 Add SHAP-based interaction features (optional, saved with seed control to keep determinism).

3.5 EDA Reproduction

• Plot feature importance bars; partial dependence for top 4 features.

3.6 Model Training & Evaluation

3.6.2 Model 1: XGBoostRegressor

- Hyperparameters: n_estimators=800, max_depth=8, learning_rate=0.05, subsample=0.8, colsample_bytree=0.8, reg_lambda=1.0, random_state=42
- Evaluation: MAE/RMSE on valid/test; calibration curve if probabilistic output is used (for classification variants).
- **3.6.3 Model 2:** LightGBMRegressor (optional) with comparable params.
- **3.6.4 Statistical Tests:** Wilcoxon signed-rank comparing absolute errors XGB vs RF across items/stores.

4 OBJECT-ORIENTED CLASS GUIDE

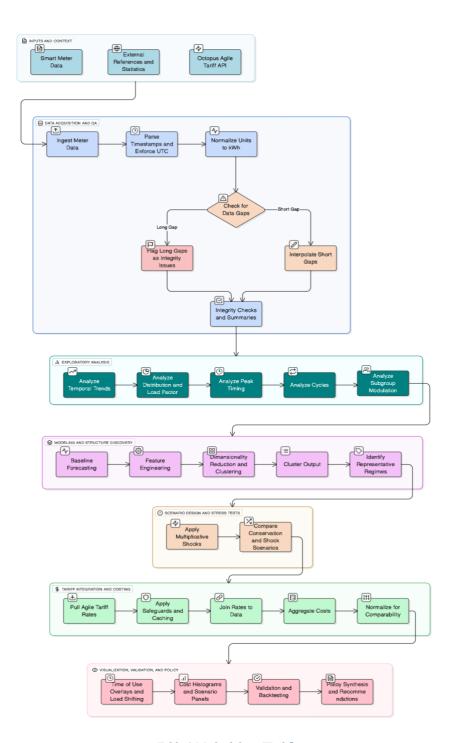


Table 1 Methodology Workflow

This supplement explains each core class in simple terms: what it does, which variables (attributes) it keeps, the main methods it exposes, and why those choices support reproducible results. Script-style utilities (e.g., set_all_seeds, savefig) are noted where they interact with classes.

Quick map of classes

~ 61	
Config utils.io	Typed container for all paths, seeds,
	and switches loaded from YAML.
RunLogger utils.io	Structured run logging (start/end,
	versions, file hashes).
DataCatalog utils.io	Central place to read/write named
	datasets with schema checks.
PlotRegistry utils.plotti	ng Standardises figure saving and
	filenames for auditability.
TariffAPIClient tariff.clien	Retrieves Agile tariff rates with
	caching, retries, and schema
	validation.
UsageDataset tariff.data	Holds half-hourly/daily usage slices
	and ensures time alignment/units.
Scenario tariff.scena	Immutable description of a tariff
	scenario (e.g., price multiplier).
ScenarioGrid tariff.scena	rio Expands
	multiple Scenario definitions into a
	design matrix.
TariffSimulator tariff.engir	Applies tariffs to usage to generate
	bills deterministically.
KPIAggregator tariff.repor	Aggregates bills into equity and
	exposure metrics.
TSPreprocessor forecasting.	preprocess Cleans and regularises time series
	(gaps, outliers, typing).
FeatureEngineerTS forecasting.	features Builds lags/rollings and
	calendar/weather features.
SarimaxModel forecasting.	models Thin wrapper
	over statsmodels SARIMAX with
	save/load.
Backtester forecasting.	evaluate Rolling-origin evaluation with
	consistent splits and metrics.

EDAReport	eda.report	Reusable routines for univariate, numeric, and bi/multivariate EDA.
DemandPreprocessor	demand.preprocess	Sklearn-style wrapper: imputation, encoding, outlier capping.
RFTrainer	demand.models	RandomForest training + calibrated metrics and persistence.
XGBTrainer	demand.models	XGBoost training + importance export and calibration.
VerificationRunner	reports.verify	Compares "reported vs reproduced" KPIs and writes a verification table.

Utility functions used by many

classes: set all seeds (seed), ensure dir (path), load cfg (path), savefig (path).

4.1 utils/io.py

Classes: Config, RunLogger, DataCatalog

4.1.1 Config

- **Purpose.** Single, typed container for parameters loaded from YAML.
- Key attributes.

```
paths: dict[str, str|Path] (e.g., raw, interim, processed, out, logs);
seed: int;
cutoffs: dict[str, str] (dates for backtests/splits);
flags: dict[str, bool] (e.g., use_cached_api, strict_schema);
models: dict (orders, horizons, tariff settings).
```

- Core methods. from yaml(path) -> Config; to dict() -> dict.
- **Rationale.** Prevents parameter drift across notebooks; ensures version-controlled, human-readable configuration.

4.1.2 RunLogger

- **Purpose.** Structured run logging for audit and replay.
- Key attributes.

```
run_id: str; cfg: Config; env: dict (package versions, OS,
CPU/GPU); log path: Path.
```

- Core methods. start(step), end(step, status), record(key, value), hash file(path) -> str.
- **Rationale.** Creates durable traces of settings and artefacts, enabling exact regeneration.

4.1.3 DataCatalog

- **Purpose.** Schema-aware loading/saving of named datasets.
- Key attributes.

```
roots: dict[str, Path]; schemas: dict[str,
dict[col,dtype]]; coerce rules: dict[str, Callable].
```

- Core methods. load(name) -> pd.DataFrame, save(name, df), exists(name) -> bool.
- **Rationale.** Enforces consistent dtypes across the four yearly half-hourly files and the daily aggregate.

4.2 utils/plotting.py

Class: PlotRegistry

4.2.1 PlotRegistry

- **Purpose.** Standardise figure export with a manifest.
- **Key attributes.** out_dir: Path; style: dict; manifest: list[dict].
- Core methods. save (fig, name, meta), list() -> pd.DataFrame.
- Rationale. Ensures reproducible, discoverable figures with metadata sidecars.

4.3 tariff/client.py

Class: TariffAPIClient

4.3.1 TariffAPIClient

- **Purpose.** Retrieve Agile prices with caching, retries, and schema checks.
- **Key attributes.** base_url: str; session: requests.Session; cache_dir: Path; timeout: float; max_retries: int; backoff: tuple; schema: dict.
- Core methods. get_rates (product, start, end) -> pd.DataFrame; validate (payload); cached (key), store (key, df).
- **Rationale.** Stabilises external data dependencies; preserves determinism via local cache and validation.

4.4 tariff/data.py

Class: UsageDataset

4.4.1 UsageDataset

- **Purpose.** Hold daily/half-hourly usage with strict typing and units.
- **Key attributes.** df: pd.DataFrame; freq: str('30T'/'D'); unit: str(Wh/kWh); segments: list[str].
- Core methods. align(other_index), to_kwh(), slice(by: dict).

• **Rationale.** Prevents silent unit mismatches and misaligned timestamps before simulation or forecasting.

4.5 tariff/scenario.py

Classes: Scenario, ScenarioGrid

4.5.1 Scenario

- **Purpose.** Immutable description of a "what-if".
- Key attributes. pm: float (price multiplier), sc_add: float (standing-charge add-on), label: str.
- Rationale. Freezes assumptions to avoid mid-run mutation.

4.5.2 ScenarioGrid

- **Purpose.** Expand multiple scenarios against usage segments.
- **Key attributes.** scenarios: list[Scenario]; cross by: list[str].
- Core methods. expand(usage) -> pd.DataFrame.
- Rationale. Clean separation of design (scenarios) from execution.

4.6 tariff/engine.py

Class: TariffSimulator

4.6.1 TariffSimulator

- **Purpose.** Deterministically compute bills.
- **Key attributes.** rates: pd.DataFrame; standing_charge: float; assumptions: dict (rounding, VAT, conversion flags); logger: RunLogger.
- Core methods. apply(usage, scenario) -> pd.DataFrame; batch(usage, grid) -> pd.DataFrame.
- Rationale. Isolates pricing logic; explicit, testable, and seed-free.

4.7 tariff/reporting.py

Class: KPIAggregator

4.7.1 KPIAggregator

- **Purpose.** Aggregate bills into interpretable metrics.
- **Key attributes.** group_by: list[str]; equity_threshold: float; indices: dict[str, Callable].
- Core methods. summarise (bills_df) -> pd.DataFrame; export (path).
- Rationale. Keeps statistical summaries separate from simulation mechanics.

4.8 forecasting/preprocess.py

Class: TSPreprocessor

4.8.1 TSPreprocessor

- **Purpose.** Clean and regularise time series.
- **Key attributes.** freq: str; gap_strategy: str; outlier_policy: dict; y_col: str; time_col: str.
- Core methods. fit(df), transform(df), fit transform(df).
- **Rationale.** Reduces variance from missingness and outliers; ensures consistent frequency.

4.9 forecasting/features.py

Class: FeatureEngineerTS

4.9.1 FeatureEngineerTS

- **Purpose.** Build lag/rolling/calendar/weather features.
- **Key attributes.** lags: list[int]; windows: list[int]; calendar: bool; weather cols: list[str].
- Core methods. apply (df) -> pd.DataFrame.
- Rationale. Centralises feature logic to avoid drift between models.

4.10 forecasting/models.py

Class: SarimaxModel

4.10.1 SarimaxModel

- Purpose. Thin wrapper around statsmodels SARIMAX.
- **Key attributes.** order: tuple; seasonal_order: tuple; exog_cols: list[str]; fitted: object.
- Core methods. fit (df), predict (h), save (path), load (path).
- Rationale. Makes model specification explicit and serialisable.

4.11 forecasting/evaluate.py

Class: Backtester

4.11.1 Backtester

• **Purpose.** Rolling-origin evaluation with fixed rules.

- **Key attributes.** horizon: int; step: int; metric_fns: dict[str, Callable]; cutoff start: str.
- Core methods. run (model, df) -> pd. DataFrame; summary (metrics_df) -> pd. DataFrame.
- Rationale. Enables comparable scores across models and folds.

4.12 eda/report.py

Class: EDAReport

4.12.1 EDAReport

- **Purpose.** Reusable routines for EDA outputs.
- **Key attributes.** registry: PlotRegistry; catalog: DataCatalog; targets/numerics/categoricals: list[str].
- Core methods. univariate yearly(df), numeric(df), bi multivariate(df).
- Rationale. Avoids one-off plotting and inconsistent filenames.

4.13 demand/preprocess.py

Class: DemandPreprocessor

4.13.1 DemandPreprocessor

- **Purpose.** Sklearn-compatible preprocessing.
- Key attributes. num_cols, cat_cols; impute_num, impute_cat; cap_strategy; encode r: OneHotEncoder.
- Core methods. fit (X, y=None), transform(X), get_feature_names().
- Rationale. Guarantees identical matrices for RF and XGB across Parts 1–2.

4.14 demand/models.py

Classes: RFTrainer, XGBTrainer

4.14.1 RFTrainer

- **Purpose.** Train and evaluate RandomForest with reproducible settings.
- **Key attributes.** params: dict; pre: DemandPreprocessor; model .
- Core methods. fit, predict, metrics, save, load.
- **Rationale.** Encapsulates preprocessing + estimator in one pipeline.

$4.14.2 \; {\tt XGBTrainer}$

- **Purpose.** Train and evaluate XGBoost with calibrated parameters.
- Key attributes. params: dict; pre: DemandPreprocessor; model .

- Core methods. fit, predict, metrics, feature_importance(), calibrate(...) (for classification variants).
- Rationale. Provides symmetric interface with RFTrainer for verification.

4.15 reports/verify.py

Class: VerificationRunner

4.15.1 VerificationRunner

- **Purpose.** Compare "reported vs reproduced" metrics.
- Key attributes. tables dir: Path; targets: list[dict]; tolerance: float.
- Core methods. build() -> pd.DataFrame; export(path).
- Rationale. Makes verification a first-class, automated step.

4.16 Shared utilities (used across files)

- set_all_seeds (seed) controls RNG for numpy and random; called early for determinism.
- ensure dir (path) safe directory creation before writes.
- load cfg(path) YAML loader returning Config.
- savefig(path) thin wrapper used by PlotRegistry for consistent export.

4.17 How files interact (at a glance)

Tariff

flow: utils.io.Config \rightarrow utils.io.DataCatalog \rightarrow tariff.data.UsageDataset \rightarrow tariff.client.TariffAPIClient \rightarrow tariff.scenario.ScenarioGrid \rightarrow tariff.engine.TariffSimulator \rightarrow tariff.reporting.KPIAggregator \rightarrow utils.plott ing.PlotRegistry.

Forecasting

flow: DataCatalog \rightarrow forecasting.preprocess.TSPreprocessor \rightarrow forecasting .features.FeatureEngineerTS \rightarrow forecasting.models.SarimaxModel \rightarrow forecasting.evaluate.Backtester.

Demand modelling

flow: demand.preprocess.DemandPreprocessor → demand.models.RFTrainer /
XGBTrainer → reports.verify.VerificationRunner.

5 GUIDE TO CODE REPOSITORY

This guide describes the repository structure and the exact, deterministic steps to run the pipeline.

5.1 File Manifest

- requirements.txt Pinned Python dependencies for a Python ≥ 3.10 environment.

 A lockfile (e.g., requirements.lock) may be exported after install.
- src/01_data_cleaning.py Ingests the SERL aggregated CSVs, validates schema, coerces types, and writes cleaned Parquet datasets. Produces run logs where applicable.
- src/02_eda.py Generates descriptive analyses and exports figures and summary tables (if configured) from the cleaned datasets.
- src/03_model_training.py Trains forecasting or demand models, writes trained artefacts, and exports evaluation metrics and diagnostic plots.
- outputs/figures/ Target directory for all generated visualisations; filenames should be stable across runs.
- outputs/models/ Persisted estimators or pipelines (e.g., .pkl/.joblib) for verification and reuse.

Recommended but optional

- o configs/ YAML configurations (paths, seeds, cut-offs, model settings) to parameterise scripts.
- o outputs/tables/ Metrics and verification CSVs/Parquet files if tabular outputs are enabled.
- logs/ JSONL run logs capturing environment details, seeds, and file hashes.

5.2 Instructions for Execution

Prerequisites. Python **3.10**+ with standard build tools; write permissions for outputs/ and logs/.

1. Create and activate an isolated environment

o macOS/Linux:

- o python -m venv .venv
- o source .venv/bin/activate
- Windows (PowerShell):
- o python -m venv .venv
- o .\.venv\Scripts\Activate.ps1

2. Install dependencies

- 3. pip install --upgrade pip
- 4. pip install -r requirements.txt

(Optional) Export a lockfile:

```
pip freeze > requirements.lock
```

5. Prepare inputs

- o Place the SERL CSVs under data/raw/ with their original filenames.
- If using configuration files, confirm the paths and parameters in configs/*.yaml.

6. Run the pipeline in order

- o Cleaning:
- o python src/01 data cleaning.py [--cfg configs/run.yaml]
- o EDA:
- o python src/02_eda.py [--cfg configs/run.yaml]
- o Modelling:
- o python src/03 model_training.py [--cfg configs/run.yaml]

7. Verify expected artefacts

- o Cleaned datasets appear under data/processed/.
- o Figures are written to outputs/figures/.
- o Trained models and metric files are written to outputs/models/ and (if enabled) outputs/tables/.
- Run logs (if enabled) are recorded under logs/ with the active configuration snapshot.

Determinism safeguards. Scripts set a master seed, record package versions, and hash inputs. Where external rates or APIs are used, caching and schema validation are enabled so regenerated runs match prior outputs.

6 CONCLUSION

6.1 Statement of Completeness

The repository structure, dependency manifest, and ordered execution steps documented above are sufficient to regenerate all figures, models, and evaluation artefacts used in the dissertation, assuming access to the source data. The empirical analyses rely on the **Smart Energy Research Lab (SERL) Statistical Data, 2020–2023**, obtained under safeguarded access from the UK Data Service (**Study 8963; DOI: 10.5255/UKDA-SN-8963-2**). With inputs placed as specified and the environment pinned, results are reproducible across systems consistent with the stated hardware and software prerequisites.

6.2 Contact for Clarification

For assistance with execution, configuration, or verification details, please include the run ID, the active configuration file (if used), and the exact script invocation. Contact details (name, institutional email, supervisory reference) should be listed here in line with departmental policy.